

**DEEP LEARNING  
PROJECT REPORT**

**School of Computer Science and Engineering**

**By**

**22BCE7775**

**Meghana Naidu T**

**2024 -2025**

## TABLE OF CONTENTS

| Chapter No. | Title  | Page No. |
|-------------|--|----------|
|             | Abstract   | 3        |
| 1           | Introduction<br>1.1 Objective<br>1.2 Scope & Motivation<br>1.3 Organization of the Report<br>1.4 Target Audience   | 4 - 8    |
| 2           | Literature Survey<br>2.1 Summarization Approaches<br>2.2 Comparison Between Extractive and Abstractive Forms of Summarization Methods<br>2.3 Specific Challenges and Solutions in Text Summarization<br>2.4 Evaluation Metrics<br>2.5 Applications of Text Summarization<br>2.6 Future Directions and the Conclusion | 9 - 22   |
| 3           | Hardware and Software Requirements<br>3.1 Hardware Specifications<br>3.2 Software Requirements   | 23, 24   |
| 4           | Proposed Methodology<br>4.1 Data Setup<br>4.2 Proposed Methodology Framework<br>4.3 Model Implementation<br>4.4 Model Training and Evaluation<br>4.5 Training Strategy<br>4.6 Evaluation Metrics<br>4.7 Comparative Analysis<br>4.8 Hyperparameter Tuning<br>4.9 Model Deployment Possibilities                      | 25 - 37  |
| 5           | Results and Discussions<br>5.1 Quantitative Evaluation<br>5.2 Qualitative Evaluation<br>5.3 Comparison of Different Models<br>5.4 Error Analysis<br>5.5 Summary of Findings  | 38 - 42  |
| 6           | Conclusion   | 42       |
| 7           | Future Work  | 42       |
| 8           | References   | 43       |

## **ABSTRACT**

The project centers on devising deep learning procedures for automatic summarization of news articles derived from text using Long Short-Term Memory (LSTM) networks. Three setups are pitted against each other to determine which produces better results: a baseline LSTM model with large dataset, a model using a deeper 2-layers LSTM architecture with smaller dataset, and one trained on a considerably smaller dataset (one which is 1000 news articles). The objectives are to produce a summary that is concise and coherent, with the main aspects of the original text intact.

The project adopts a sequence-to-sequence model including encoder-decoder structures with attention mechanisms to achieve further context retention in summary production. Preprocessed news datasets are used in training the models, with text cleaning, tokenization, padding, and vocabulary optimization as steps in the processing. Performance measurement Int by ROUGE metrics, qualitative inspection of the generated summaries, and training loss analysis.

The results reveal that complex architectures and finally larger datasets afford more context-aware and fluent summaries than simple models or small datasets. For bigger datasets, these advantages come at the cost of longer training times and computation time. The project demonstrates the pair of LSTM models for abstractive summarization tasks, while at the same time disclosing the trade-offs with the size of data collection and depth of architectures.

# 1. INTRODUCTION

When the amount of digital content, especially by online news websites, bloggers, and social networks, increases rapidly, users end up having to deal with large daily volumes of text information. But with the great potential in this exposition to find knowledge also comes the challenge of finding, extracting, and comprehending relevant content efficiently. Automated text summarization is one of the serious solutions presented in such a situation, allowing users to understand the content of long texts without actually reading through them.

Text summarization is a sub-field of Natural Language Processing (NLP) devoted to reducing large amounts of text to smaller portions that still reflect the key information and the meaning of original messages. The important two types of text summarization are: Extractive and Abstractive. Extractive summarization, in general terms, is the selection and merging of certain sentences or phrases from the source text for generation of a summary, while in abstract summarization, one attempts to rephrase and produce new sentences that carry the same kind of information. Abstractive summarization is more difficult, but it has more flexibility and readability and is more like human summarization.

The current state of the art in deep learning comprises the application of Recurrent Neural Networks (RNNs) and, their further development Long Short-Term Memory (LSTM) networks, which has broadened the spectrum of the abilities of machines in understanding and generating human language. LSTM networks have a particular emphasis upon learning long-range dependencies from sequential data and hence are some of the most important candidates for increasingly advanced NLP tasks such as machine translation, speech recognition, and summarization.

The current project capitalizes on the process of the LSTM-based deep learning models in performing abstractive text summarization of the news story. An attention mechanism Seq2Seq model has been developed to enhance the model's capability to focus on critical aspects of the input while generating summaries. Several configurations of the model have been constructed and evaluated against one another to understand the variation in the architecture (such as number of LSTM layers) and dataset size effects on summary quality and coherence.

News articles are chosen primarily because they have real-life relevance and are ill-structured, making them amenable to the process of summarization. In automating the summarization of news, the project intends to bring time-saving and better comprehension benefits for readers, researchers, and analysts.

This document provides a comprehensive report on summarization with LSTM models, which takes into account the following elements: data preprocessing, model architecture, training procedure, evaluation metrics, and result analysis. Thus, through this comparative analysis, the project identifies among various types of LSTM configurations the advantages and disadvantages and demonstrates the way of deep learning-thinking in changing the way I read textual content.

## 1.1 Objective

Deep learning is the main aspect of abstractive text summarization of the news articles. This project is set out to model, implement, and test the deep learning models for abstractive text summarization based on the LSTM sequence-to-sequence architecture. The project appears determined to test the

influence of variations in size of LSTM network configurations and train data on performance, accuracy, and coherence of summaries.

Some of the major aims of the project include:

- To set a baseline text summarization model with a simple encoder-decoder 2-layer LSTM architecture to learn to generate short summaries from raw news articles.
- To further improve the performance of the model by experimenting with deeper network structures, particularly a 3-layer LSTM configuration, and observe the disturbance in context understanding and language generation.
- To assess the influence of the amount of training data by comparing model performance when trained over a small set vs. that trained over a large set dataset.
- To implement attention mechanisms in the LSTM architecture-based models to allow the decoder to focus on specific areas of interest in the input while producing summaries that improve the semantic relevance of the output.
- To subject the output summaries to qualitative evaluation (coherence and readability) and quantitative analysis with ROUGE scores and training losses.
- To analyze trade-offs and best practices for the design of LSTM-based text summarization models in respect of their depth, size of training data, and computational power.
- These goals intend to facilitate the establishment of a working and scalable news summarization system to enhance content consumption and information access in real life.

## **1.2 Scope & Motivation**

### **Scope**

The project intends to design and evaluate deep learning models for abstractive text summarization, specifically concerning news articles. The underlying model employs Long Short-Term Memory (LSTM) networks based on their strength of working with sequential data and remembering long-term dependencies in text.

To evaluate the performance of LSTM-based summarization, various types of the encoder-decoder architecture are built, which encompasses a baseline single-layer model, a more complex two-layer LSTM model, and a model that sees training on an extended dataset of 40000 news articles. Comparing in this way helps us study the effect of model depth and dataset size on performance.

Attention-enhanced sequence-to-sequence (Seq2Seq) learning techniques are employed by the project. Attention enables the model to focus on different parts of the input sequence while generating summaries, fostering context relevance and overall quality.

Preprocessing is critical in maintaining data quality and uniformity. The preprocessing of these news articles includes normalizing text, removing stop words, tokenizing, and padding. This procedure

emphasizes input data being clean, standardized, and ready for an efficient modeling process.

Modeling performances are assessed through qualitative and quantitative analyses. Quantitative evaluation checks the ROUGE (Recall-Oriented Understudy for Gisting Evaluation) metrics for the accuracy of generated summaries versus the reference summaries. The variables for training loss and validation loss are analyzed for their behavior over time. Also, human assessment of summaries focuses on evaluating fluency, coherence, and relevance.

A major component of the project is juxtaposing how different factors influence summarization performance, such as depth of the LSTM network, size and diversity of the dataset, and the presence or absence of attending mechanisms. Such comparisons are essential for understanding which architectural and data-oriented decisions enhance summarization performance.

Though targeted at news content summarization, the methods and insights from this work can be applied in a wider context, including legal documents, scientific articles, blogs, etc. In terms of potential applications of this project, its concise history does emphasize supervised learning based on Ill-labeled datasets. It is exciting to think of investigating other kinds of models, such as transformer-based ones, like BERT or GPT. However, currently, that is deemed outside the program scope.

## **Motivation**

An overload of textual data online deals an increasing burden to an individual today. Online news sites publish tons of articles daily, which makes it difficult for a reader to remain updated without investing considerable time in this effort. Summarization is the automated means of relieving such a burden through the production of concise and important representations of lengthy documents.

Thus, the factors inspiring this work are the necessity for faster information consumption and summarization across journalism and education, as Ill as research; the potential of deep learning-LSTM networks in comprehending and producing human-like text-for summarization in a more natural and accurate way; an interest in exploring the power of sequence models to learn the semantic and syntactic structures of texts and push their boundaries through experimentation with architectures and datasets; and ultimately further development of the system for abstractive summarization, which in its still-hard NLP task depends much on language generation rather than sentence extraction.

This project intends to promote information accessibility by automating news content summarization and laying the groundwork for practically meaningful intervention approaches for intelligent content filtering and summarization systems.

## **1.3 Organization of the Report**

The report is organized into several clearly defined chapters that together introduce the rationale, methodology, result implementation, and evaluation of LSTM-based text summarization models on news articles. Each chapter is written clearly for a logical and synthetic progression in the mind of the reader. The organization is as follows:

- i. Introduction

Gives a broad overview of the problem of information overload and the need for summarizing text. It gives the aims, scope, motivation for the project, and the structure of the report.

ii. Literature Survey

Surveys recent research and methods in an area of text summarization. The emphasis is going to be on the transition from classical extractive methods to contemporary neural network-based abstractive methods with specific mention to contributions from sequence models such as LSTMs and attention mechanisms.

iii. Methodology

Describes the methodology adopted in the project. This includes the gathering of data and pre-processing, structure of the LSTM-based models (baseline, 2-layer, and extended data variant), and explanation of the encoder-decoder method with attention.

iv. Implementation

Describes a chapter-wise guide to model building in Python and deep learning frameworks like TensorFlow and Keras. Discusses hyperparameter tuning, model training, and issues with overfitting or training.

v. Results and Evaluation

A chapter giving performance analysis of all models. The comparison includes ROUGE scores, training vs. validation loss curves, and sample summaries assessing output quality and relevance.

vi. Conclusion and Future Work

An account of summarizing the notable results from the project and an evaluation of the varying performances for the models. Also, suggestions for further improvements are given, like exploring transformer-based models or real-time summarization systems.

vii. References

Lists all papers, tools, datasets, and libraries that are cited throughout the development of the project and in writing the report.

## 1.4 Target Audience

This is a broad project directed at both the academic as well as industry communities of people interested in deep learning and applications of natural language processing. The primary beneficiaries of this research include:

a) NLP Researchers

As an introduction to the effects of LSTM-based encoder-decoder models and attention mechanisms on abstractive summarization tasks, this study will be useful to academic and practical people engaged in the development of natural language processing.

b) Journalists and Media Analysts

Increasingly adopting automated content curation, news agencies might benefit from insights on how to develop succinct and meaningful summaries from long news articles, thus saving time while increasing efficiency.

c) Developers Creating Summarization Tools

The integration and design scheme of an evaluation technique discussed in this documentation could serve as the underpinning platform for the development or advancement of production-grade applications for software developers and engineers constrained to design AI-augmented summarization systems.

d) Researchers in Deep Learning and Computational Linguistics

Deep learning methods researchers and students for language generation would find this research most valuable because they will understand real-world challenges while designing models and choices in sequence-to-sequence learning and abstractive summarization.

This is intended to provide a bridge between these groups through closing the gap between theoretical research and practical deployment of deep learning models in text summarization.



## 2. LITERATURE REVIEW

In this present age of enormous information overload, text summarization is an important tool in natural language processing (NLP) that has the task of compressing large volumes of text to shorter versions, retaining most of the original content. The essential motivating factor for summarization is to make information more accessible, reduce reading time and facilitate quicker decision-making in various areas such as journalism, law, scientific research, e-commerce, and software development.

Broadly, summarization can be subdivided into two main parts:

- i. Extractive summarization- involves selection and aggregation of the most relevant sentences or phrases from the original documentation.
- ii. Abstractive summarization means the generation of new sentences that can convey core ideas, thus producing more coherent and human models of summaries.

Summarization was earlier based on heuristic techniques such as frequency scoring (like TF-IDF), position importance, and graph-based algorithms like TextRank and LexRank. Although computationally efficient, they have serious drawbacks in terms of semantic understanding most times and how coherent or non-redundant the produced summary is.

With deep learning came a big changeover. The introduction of sequence-to-sequence (Seq2Seq) models-based on Recurrent Neural Networks (RNNs), especially Long Short-Term Memory (LSTM) networks, has made the system learn better dependencies and context within texts. When combined with attention mechanisms, these models translated a vital turning point for abstractive summarization, allowing attention from models on very salient aspects of the input generation.

Pointer-generate networks, transformer architectures (such as BERT, GPT, and DistilBERT), generative adversarial networks (GANs), and the likes did not exhaust on factual inconsistency, repetition, poor generalization, and other such issues which summary systems are known to encounter. Hybrid-incorporating both extractive and abstractive techniques-has also gained much attention for its balance between coherence and informativeness.

The papers reviewed in this study-from GAN-based abstractive models to deep clustering for legal document summarization and BERT-driven extractive approaches-are evidence of this evolving trend. They illustrate the shift from rule-based extraction through deep learning to semantically rich and context-aware generation.

Thus, this review will expand to appraise in detail these modern approaches, their methodologies, strengths, limitations, and contributions toward automatic text summarization.

### 2.1 Summarization Approaches

The methods for text summarization seem to present changes as the adoption of textual information grows and users expect efficient retrieval for information. These techniques are further categorized into two broad categories: extractive and abstractive summarization. Both approaches are aimed at shortening long documents; however, their methods of doing so vary drastically, including their architectures and applications.

#### 2.1.1 Extractive Summarization

Extractive summarization does not paraphrase or generate new content. They only extract the most representative sentences from the document about the ideas it presents. This method is simpler to put into practice and is less prone to grammatical errors since it does not require natural language generation.

## 1. Traditional Methods

The classical extractive methods relied heavily on statistical and heuristic-based techniques. These approaches are computationally efficient and interpretable, but often lack depth in understanding the semantic content of a document.

*Classification of methods on the basis of frequency:* Among the earliest methodologies, Term Frequency-Inverse Document Frequency (TF-IDF) is applied to measure the importance of words about their occurrence within a document and in an entire corpus. Subsequently, the sentences are assigned scores based on the sum of the TF-IDF scores of the words contained in them. Simple and fast though they may be, due to simplicity such methods quite often lack sensitivity to the context, semantics, and redundancy among sentences that they summarize.

*Position-Based Scoring:* Some models score more highly for sentences that come first based on the notion that journalistic writing will tend to place important information at the beginning. This heuristic may, however, be domain-dependent and unreliable across a plethora of textual formats.

*Graph-Based Techniques:* TextRank, whose principle is based on Google PageRank algorithm, and LexRank model sentences as nodes in a graph, semantically linked by edges (usually cosine similarity of TFIDF vector). Summaries are created by selecting top-ranked nodes (sentences). These unsupervised models identify central sentences brilliantly—yet they struggle with redundancy and don't have any firm semantic ground to stand upon.

These classical methods laid the foundation for extractive approaches to summarization but fell short in terms of contextual and syntactic deficits. As documents became more intricate, especially in the legal, medical, or technical fields, these methods could not handle such complexity.

## 2. Deep Learning-Based Techniques

With the advent of deep learning and contextual embeddings, something interesting happened in extractive summarization. These models capture not just statistical properties of texts but also the semantic, syntactic, and contextual features of sentences and documents.

*Contextual Embedding with BERT:* Use of Bidirectional Encoder Representations from Transformers (BERT) is perhaps the most far-reaching innovation into extractive summarization. When pretraining BERT models is followed by fine-tuning for downstream summarization tasks, they encode sentences into dense vectorial embeddings that include a lot of deep contextual information.

In "A sentence is known by the company it keeps: Improving Legal Document Summarization Using Deep Clustering" (Jain et al., 2024), a Legal BERT domain-specific model was applied for scoring the relevance of sentences in legal documents. The embeddings produced are then classified via an MLP into summary-relevant and irrelevant categories. Sentence scoring captures local context, whereas the paper introduces deep clustering for addressing global document comprehension, thus enhancing coherence and coverage.

*Graph Neural Networks and Hybrid Architecture:* Newer approaches integrate the BERT-based sentence embeddings into graph structures or clustering algorithms. Certain models cluster similar sentences and select representatives from each cluster to ensure topic diversity and reduce redundancy.

*Integration of Reinforcement Learning:* The paper "A generative adversarial network-based extractive text summarization using transductive and reinforcement learning" (Yang et al., 2023) describes how reinforcement learning is used to fine-tune the discriminator of a GAN-based extractive model. Here, sentence selection is predicated on a sequential decision-making pattern namely maximizing a trade-off between precision and recall. Reinforcement learning also solves the problem of class imbalance, which is a standard problem in most extractive datasets.

*Transductive Learning:* In the same paper, transductive long short-term memory (TLSTM) is introduced into the GAN architecture. TLSTM prioritizes test-time samples, attributing higher weights to training samples that are similar to test samples. This way, it enhances the generalization of the model and tailoring of summaries to the specific document context.

*Efficiency via DistilBERT:* To resolve the heavier resources requirements of BERT, DistilBERT offered a lighter alternative. While it may be marginally less expressive than its less-distilled counterpart, it increases inference tempo and memory efficiency to a considerably greater extent-in other words, it is more well-suited for real-time applications.

### **Advantages of Deep Learning-Based Extractive Approaches:**

- Capture deep contextual information.
- Handle synonyms, paraphrases, and domain-specific terms much better.
- Help lend toward supervised learning with the use of labeled datasets.

### **Disadvantages:**

- Still limited to verbatim sentence extraction-may admit verbose content or contextually disjointed content.
- Can produce incoherent summaries when not coupled with redundancy reduction techniques.
- Require huge labeled datasets for training and huge computational resources.

### **2.1.2 Abstractive Summarization**

The abstractive summarization is different from extractive one that is supposed to put together phrases in order to make new sentences that represent the meaning of the source text. This is just the way humans summarize by understanding first what it is all about then paraphrasing or rephrasing it into their own words. Hence, more often it's found that abstractions are fluent, coherent, and non-redundant.

Practically, it is very complex. There are very many hurdles of generation like grammatical, factual, and semantic accuracy and relevance which are common to all forms of natural language generation.

Over the years, abstract summarization research has evolved with the growth of sequence modeling, attention mechanisms, and deep learning architectures.

## 1. Sequence-to-Sequence (Seq2Seq) Models

Seq2Seq: This is the basic template architecture of abstractive summarization, comprising:

Encoder: This translates input text into a fixed, length vector representation.

Decoder: This takes that encoded representation and generates a summary, word by word.

This was first popularized in the area of machine translation but subsequently was adopted for summarization. both the encoder and decoder are often such models as LSTM or GRU. As highlighted in "Abstractive Text Summarization using Adversarial Learning and Deep Neural Networks" (Tank & Thakkar, 2024) -

These Seq2Seq models enabled end-to-end training pipelines and facilitated the learning of semantic alignment between input documents and summaries by such models. However, the constructs bore a remarkable departure from rule-based summarizers, with limitations including- a fixed-length vector bottleneck, ineffectiveness in Long-range dependencies, exposure bias due to teacher forcing during training (i.e., the model sees ground truth during training, but not during an inference operation).

## 2. Attention Mechanisms

To tackle the fixed-length vector problem that Seq2Seq models are widely known for, attention mechanisms were introduced. The attention mechanism allows the decoder to dynamically focus on different parts of the input sequence at every time step. The model can selectively focus on the more relevant parts of the source text while generating each word of summary.

This was first described by Ashish Vaswani and others and has since been developed further by Chopra et al., 2016, and See et al., 2017; it paved the theoretical ground for many modern summarization systems. In Tank and Thakkar's paper the generator of their adversarial network utilizes such attention-driven encoder-decoder architectures that much more effectively align source sentences with generated summaries. Pan-Gao mentioned that, however, vanilla attention mechanisms do still suffer from factual inconsistencies, OOV (out-of-vocabulary) terms, and content repetitions.

## 3. Pointer-Generator Networks

To boost the quality of generation even further, See et al. (2017) proposed a model called the Pointer-Generator Network (PGN), a hybrid model combining abstractive generation with extractive copying.

Key design features:

*Pointer mechanism:* This allows the model to copy words directly from the source text—especially useful for names, numbers, or OOV terms.

*Generator mechanism:* This allows the model to generate new words from the vocabulary.

*Switching network:* This is a trainable mechanism determining whether to copy or generate at each step.

This innovation contributed greatly to improving accuracy with respect to facts, OOV words, and the coherence of the summary. As mentioned in the survey of Kumar et al. (2024), coverage mechanisms are helpful for PGNs not just via reducing word-level and sentence-level repetition but also for

increasing the likelihood that content previously attended to will not be repeated in future steps.

#### **4. Transformer Models**

The transformer model (introduced in Vaswani et al., 2017) swept through NLP by entirely avoiding any recurrence and leveraging a multi-head self-attention mechanism.

Transformers are completely parallelizable, thus speeding up training, capable of capturing long-term dependencies better as the self-attention can attend to all the input tokens, achievable and flexible depending on single-document or multi-document summarization.

Some of the transformer-based models which gained significance in the context of abstractive summarization include:

- BERT (Bidirectional Encoder Representations from Transformers): Mainly useful for extractive tasks; however, BART and PEGASUS are variants adapted from the transformer for generation.
- BART (Bidirectional and Auto-Regressive Transformers): BART was pretrained using a denoising autoencoder objective, hence it is very capable of generating high-quality abstractive summaries.
- T5 (Text-To-Text Transfer Transformer): Considered all tasks of NLP as a text generation task, providing a consistent framework for summarization and beyond.

In this work titled "Product Backlog Recommendations Using Sentiment Analysis With Topic Modelling and Text Summarization in the XYZ E-Wallet Application" (Susanto & Mauritsius, 2025), BART is available for summarizing customer feedback in both the positive and negative sentiments. Owing to its emotional inferencing and factual retaining characteristics while paraphrasing, it is suitable for business as well as consumer summarization.

#### **5. GAN-Based Approaches in Abstractive Summarization**

A standard GAN framework consists of two components: the Generator (G) which tries to create summaries that would look like they were written by a person, and a Discriminator (D) that tries to distinguish between human-created and machine-generated summaries. In the case of the DeepTextSummGAN model proposed by Tank and Thakkar, the generator is taken as the Seq2Seq model emboldened with attention mechanisms, while the discriminator is a CNN-based classifier. The generator is optimized using reinforcement learning, more specifically a policy gradient approach, for producing summaries that fool the discriminator thereby enhancing the ability of producing good human-like outputs. The merits of this technique are that it reduces exposure bias; allows for direct optimization of evaluation metrics such as ROUGE through reward feedback; and results in coherent, human-like summaries. But there are lots of complications when it comes to the application of GANs to NLP due to the discrete nature of text, which makes optimization harder.

Abstractive summarization undoubtedly has immense benefits but also some pronounced drawbacks. It can be used to create fully formed sentences that sound natural, for example, and quickly reduces repetition as compared with extractive methods. This makes it even suitable for highly creative summary tasks like story, conversation, or review summaries. This capability also allows for abstraction and condensation from hard-to-approach large or unstructured input. These benefits are, however, overshadowed by the various challenges copy-pasting their way into this space, such as the

likelihood of producing factual inconsistencies and hallucinations. Additionally, evaluating the quality of an abstract summary remains a challenge since traditional measures such as ROUGE do not capture the deep semantic fidelity or coherence of the summary.

## 2.2 Comparison Between Extractive and Abstractive Forms of Summarization Methods

While speaking of the growing field of automatic text summarization, both types have adapted with time-from rule-based and statistical towards deep learning and hybrid architectures. Every summarization paradigm has its own benefits and restrictions, per se influencing their choice based on domain, availability of data, and performance metrics.

### 1. Methodology Core Principle:

| Aspect                     | Extractive Summarization                                       | Abstractive Summarization  |
|----------------------------|--|--|
| <b>Core Principle</b>      | Selects and concatenates important sentences from the source   | Generates new sentences capturing the core meaning of the source |
| <b>Data Requirement</b>    | Can work with limited or unlabeled data                        | Requires large labeled datasets for training                     |
| <b>Generation</b>          | Copy-based (verbatim from text)                                | Natural Language Generation (NLG)-based                          |
| <b>Techniques</b>          | TF-IDF, TextRank, BERT embeddings, clustering, GANs            | Seq2Seq, Attention, PGN, Transformers, GANs with RL              |
| <b>Domain Adaptability</b> | Easily adaptable with minimal training (e.g., legal summaries) | Requires domain-specific fine-tuning for best results            |

### 2. Performance and Accuracy

| Dimension                    | Extractive Methods                                | Abstractive Methods   |
|------------------------------|---|---|
| <b>Coherence</b>             | May lack logical flow between selected sentences  | Often more fluent and cohesive                              |
| <b>Factual Consistency</b>   | High (as sentences are copied directly)           | May hallucinate or misrepresent facts                       |
| <b>Redundancy</b>            | Often includes repetitive content                 | More efficient at compressing and rephrasing                |
| <b>Handling of OOV Terms</b> | Well-handled if term is included in original text | Struggles unless copy mechanisms are integrated (e.g., PGN) |
| <b>(ROUGE Scores)</b>        | Generally strong with conservative coverage       | High variability depending on architecture and training     |

For example, within the DeepTextSummGAN model, the abstract summaries produced superior ROUGE-1 (41.58) and ROUGE-L (38.96) scores over all other models previously, which indicates that performance gap occurs with the adversarial learning. On the other hand, for legal documents, DCESumm-an extractive approach-improved the quality of summaries in terms of relevance and the F1 score on BillSum and FIRE datasets, hence showing evidence of its effectiveness in the domain even when sentence generation is not available.

## 2. Strengths and Ideal Use Cases

Despite their essential similarities, extractive and abstractive summarization differ relatively conspicuously in many highly relevant details. Simplified and fast, extractive summarization is generally faster and quite computationally lightweight, while on the other hand, the very fact that, in general, more sophisticated machine techniques based on natural language generation and heavy training are being employed for abstractive summarization makes it almost always more time-consuming.

Extractive summarization is suitable for text where phrasal integrity is a must, as in structured formal documents like laws, medicine, and technical texts. With this, abstraction stands as a good option for narrative resources, reviews, social media posts, and where the power of rephrasing and interpretation adds value.

In terms of reusability, extractive models generally allow for far-reaching generalizations for which little domain-specific fine-tuning is required. Abstractive methods, however, usually require some significant domain-specific training in order to guarantee decent quality translations. Departure from reliable text interpretation points to the complication associated with the latter garbage in and out scenario. In the case of extractive summaries, the axiomaticity of analysis is simple because one can pinpoint the exact sentences from which the statement was picked, whereas abstractive summaries generate new phrases that are hard to associate with the original text.

Last, with regard to fluency and diversity, are quite limited in extractive summarization by the original formulation of the source text, whereas reading and profundity is essentially what we multicasting.

In the case of the XYZ E-Wallet study, BART, an abstractive summarizer, did well in synthesis positive and negative reviews at separated collections or summaries that fit into structured forces of customer feedback loops. The GAN-TLSTM extractive model, on the contrary, is more concerned about the accuracy of sentence inclusion and coherence with respect to DistilBERT evidence, which makes it very suited for "efficiency-critical" tasks.

## 4. Challenges

Applications of extractive and abstractive summarization methods indeed pose different problems to the user. For instance, extractive summarization may at times pull out what is being said away from the context and lead to semantically erroneous summaries that are otherwise factually correct. Abstractive summarization, on the other hand, happens to be better at contextual paraphrasing, thus rendering content more closely matching the meaning of the source text.

The training requirement is also completely different. Therefore, an extractive summary may not depend much on data use where for example a small amount of training may do well. In contrast, for an abstractive summary, much labeled data is needed together with huge expense on a high computational machine for it to run well for high-quality results.

Semantic comprehension stands as another aspect in which the two differ. Extractive models are limited in semantic comprehension due to the fact that they simply select parts of the input verbatim. Abstractive models, on the other hand, possess a stronger grasp of semantics that allows them to change the ideas in new ways but are also prone to hallucination in which the content generated may not be factually supported by the source.

Last but not least, evaluation complexity is higher for abstractive summarization. Simple overlap-based metrics such as ROUGE that measure the number of words or phrases that the summary shares with the reference text. On the other hand, the assessment of the abstracting techniques requires greater nuance with respect to semantic accuracy and factual correctness, making the process much more complicated.

Both extractive and abstractive summarization techniques bring essential value depending on the goal, domain, and available resources. Extractive methods, especially those enhanced by deep learning and clustering, are very effective in structured domains like legal and scientific texts. They are composed of Seq2Seq, attention, and Transformers. In contrast, abstractive methods do better in creating more human-like, coherent, and concise summaries for creative or narrative applications.

## **2.3 Specific Challenges and Solutions in Text Summarization**

Although great strides have been made in the development of text summarization techniques, there still exist some challenges which remain unsolved and inherently affect the robustness and generalizability of the system. Majorly, these limit the use of the techniques in terms of types of documents to summarize, redundancy, and, to some extent, computational cost. Most recent research and the selected papers reviewed cover most of these problems that have been addressed in innovative ways.

### **1. Addressing Lengthy Documents**

*Challenge:*

Documents have been very lengthy and currently difficult to summarize by many earlier models of Summarization such as Seq2Seq- or LSTM-based models because of limitations in memory; sequence dependency, and; attention span. That is because RNNs and vanilla attention mechanisms always lose the context or either fail to attend to the beginning or end of the document effectively; more so in the fields of law, medicine, or science when they have multiple pages of document.

*Solution:*

Deep Clustering Enhanced Summarization (DCESumm), as proposed by Jain et al. (2024), tackles this by combining:

- i. Sentence-level scoring using Legal BERT embeddings.
- ii. Unsupervised clustering to capture document-wide structure.

Instead of pure local sentence importance based scoring, DCESumm employs deep embedded clustering for grouping semantically similar sentences and increasing their score value according to context of cluster. This ensures that long documents are parsed accurately and summarized with a global consistency. Besides, models like Transformers (e.g. BART, PEGASUS), by their self-attention mechanism, remember a longer sequence. They all still pay inefficiencies.

### **2. Ensuring Coherence and Relevance in Summary**

*Challenge:*

One of the persistent difficulties in both extractive and abstractive summarization is ensuring that the generated summary is semantically coherent and contextually relevant. Extractive summaries may consist of disconnected sentences, while abstractive summaries can drift off-topic or "hallucinate" facts that are not present in the original text.

*Solutions:*



#### a. Adversarial Learning with GANs

Tank & Thakkar (2024) in their paper introduced the DeepTextSummGAN, an adversarial framework in which:

- i. The generator produces summaries.
- ii. The discriminator tries to learn the differences between real (human-created) summaries and machine-generated ones.

This adversarial interaction encourages the generator to imitate human-like writing, increasing the fluency, coherence, and factual alignment of the summaries. Train the model by reinforcement learning to align the generator's objective with the evaluation metrics such as ROUGE.

#### b. Pointer-Generator Networks and Coverage Mechanisms

These point to the processes described in the review by Kumar et al. (2024), in which pointer-generator networks (PGNs) balance copying exact phrases (ensuring factuality) and generating novel ones (ensuring coherence). PGNs use coverage vectors to track what content was attended previously, thus reducing redundancy and ensuring better flow and coverage.

### 3. Reducing Redundancy

#### *Challenge:*

The majority of early summarization models, especially extractive, share the glaring fault of selecting sentences with overlapping content, thereby leading to verbose and repetitive summaries. This is also true with some exceptions for abstractive models, which may repeat certain phrases or sentences if not duly constrained.

#### *Solutions:*

##### a. Deep Embedded Clustering (DCESumm)

In short, a reduction of redundancy in DCESumm means:

- i. Grouping semantically similar sentences by deep clustering.
- ii. Selecting diverse sentences from different clusters.

This guarantees the summary addresses different key ideas instead of reiterating the same concepts in diverse forms. Clustering helps scoring sentences that allows for balancing between each single sentence's importance and overall document diversity to avoid redundancy in domains like legal writing.

##### b. Summarization GANs with the Discriminators

The discriminator feedback loop penalizes repeated content while rewarding novel, coherent outputs in GAN-based models. This adversarial training discourages the generator from recycling phrases and structures, especially when doing so does not contribute to the summary's overall informativeness.

##### c. Attention and Coverage Mechanisms

The critical role that coverage mechanisms in attention (in PGNs or variants of Transformers) play in recording history of attention means repetition in different forms is hardly done, Kumar et al. (2024).

## 2.4 Evaluation Metrics

### a. ROUGE

This metric is a general sceptic referring majorly to automatic evaluations of summaries with the

generation of raw summary output against deep-human recommendations. The emphasis is on n-gram, word sequence, and word pair overlap between the generated and reference summaries. Major variants include ROUGE-1, ROUGE-2, and ROUGE-L.

- i. ROUGE-1 takes into consideration the individual words in the generated summaries that can be found in reference summaries.
- ii. ROUGE-2 refers to bigram overlaps, catching short matches by phrases.
- iii. ROUGE-L concerns itself with the longest common subsequence of LCS between the generated summary and reference summary, thereby capturing sentence-level fluency and structure.

An important majority of the papers including: Tank & Thakkar, 2024 ;Yang et al., 2023 ; Jain et al., 2024 make use of ROUGE-1 and ROUGE-L as their primary evaluation indices. For example, DeepTextSummGAN, which reported ROUGE-1 = 41.58 and ROUGE-L = 38.96, has shown an advantage over GAN-based models for the CNN/DailyMail dataset.

On the other hand, the GAN-based extractive model with TLSTM and RL reported ROUGE-1 = 52.45 showing good enhancement with respect to precision and informativeness.

## **b. Emerging Semantic Metrics**

In recent literature such as Kumar et al., 2024

the following semantic evaluation methods are put forward as alternatives to ROUGE:

- i. BERTScore primarily uses contextual embeddings from BERT to measure inference of the generated summary against the reference summary; making it less sensitive to lexical variation.
- ii. Motor Score minimizes "effort" to align the summary against the reference.
- iii. Meteor and BLEU mostly rely on technicalities of machine translation, although they are sometimes used for evaluating summarization mainly from the perspective of precision, recall, and synonymy.

## **c. Benchmark Datasets**

They are instrumental in specifying the inner workings of summarization models, primarily understanding their generalization capabilities and comparative analysis with the best systems in the field. These datasets usually consist of source documents and human-written summaries for consistent model evaluation.

### **i. CNN/DailyMail**

The CNN/DailyMail dataset is predominantly constructed from news articles and headlines and is thus suitable for multi-sentence document summarization. The purpose behind its construction is to train and test models for abstractive summarization. However, it consists of some 300K document-summary pairs for training, which makes the dataset an industry standard, possessing ample structure and enough vocabulary to build and compare many advanced models. Areas of research with this dataset include the training of DeepTextSummGAN and GAN-RL-TLSTM models. Source text from online news articles and bullet points for a summary fit well with models designed for brief document

summarization.

## ii. BillSum

BillSum analyzes U.S. Congressional bills and their summaries, which makes it pertinent to summarizing legal or governmental documents. The U.S. Bills subset itself has about 23,000 instances. The whole idea behind the dataset is to help the evaluation of summarization models in dealing with fairly formal and structured language. It has been cited in particular with the focus of the DCESumm paper, which addressed the enhancement of extractive summarization of legal texts using deep clustering methods. Clustering methods in conjunction with Legal BERT led to a 6% gain in the F1-score on this dataset, illustrating that domain-embedded models could enhance summarization performance.

## iii. FIRE Legal Dataset

The FIRE Legal Dataset comprises Indian legal case documents annotated with extractive reference summaries. Therefore, it can be considered a testing ground for different summarization methodologies on the legal paradigm in the Indian judiciary. BillSum and the FIRE Legal Dataset were jointly applied in the DCESumm paper, giving rise to cross-domain validation for legal summarization tasks. Inclusion of this dataset would aid in testing adaptability for models trained on legal documents from different jurisdictions, fortifying the work on more generalizable and robust summarization methodologies for the legal domain.

In conclusion, metrics and dataset analysis while it comes to measuring summary quality, ROUGE remains the apparent favorite; however, its shortcomings have led to the proliferation of semantic metrics like BERTScore and MoverScore that provide a more fine-grained assessment. Likewise, datasets such as CNN/DailyMail serve as multipurpose benchmarks for conventional summarization, while BillSum and FIRE provide the necessary complexity and specificity to shape advancement in legal and high-stakes domains.

Good summarization evaluation nowadays typically includes:

- a. Metrics (ROUGE, BERTScore).
- b. Domain-specific datasets.
- c. Evaluations done by humans (in terms of coherence, fluency, and informativeness), particularly when the ROUGE scores tend to saturate or misalign with human perception of quality.

## 2.5 Applications of Text Summarization

Text summarization is a primary natural language processing (NLP) endeavor with a multitude of real-world applications across several industries. With applications in the law, the media, business, and software engineering, summarization contributes to increased productivity through reduced cognitive load and rapid accessibility to critical information. An elaborate description centralizes application in four notable domains:

### 1. Legal Document Summarization

*Importance:*

Legal documents, such as case judgments, legislative bills, and contracts, can be very lengthy, highly technical, and require difficult wording. Summarizing such documents would give brief opportunities for lawyers, judges, policy analysts, and citizens to quickly orient themselves with the legal document

contents, aiding in faster decision-making and transparency.

*Impact:*

The DCESumm would assist legal practitioners with identifying the most relevant portions of a document, citizens in the reading of laws affecting them, and government bodies in improving access to the public policy documents.

## **2. E- Commerce**

*Importance:*

Text summarization is very significant in e-commerce: it is auto-generating product descriptions, summarizing user reviews, and highlighting product features or shortcomings. With thousands of user reviews for any product, summarization is a scalable solution for surfacing valuable insights.

*Impact:*

These summarized insights are converted into product backlog recommendations, streamlining the feature development process and improving customer satisfaction. Such an approach helps companies, such as e-wallet platforms or online marketplaces, prioritize these updates based on actual user problems.

## **3. News Summarization**

*Importance:*

News summarization is the solution to the information overload that darkens the modern reader. It enables quick consumption of breaking news, policy changes, and investigative reports. News agencies, aggregators, and even social platforms use summarization to create headline previews, daily digests, and personalized feeds.

*Impact:*

News summarization is enabling media platforms to communicate real-time concise updates to their users, enabling them to compare multiple articles, and driving features like smart assistants and news bots.

## **4. Recommendations for Product Backlog in Software Development**

*Importance:*

In agile software development, the dynamic product backlog will include the features, bug fixes, and technical tasks to be worked on in future sprints. Traditionally, a product manager or developer would curate the backlog, but the quality of the backlog could be enhanced significantly by means of real-time user feedback.

*Impact:*

- i. Enables data-driven decision-making for feature development.
- ii. Lighten the load on manual analysis teams.
- iii. Improves the customer satisfaction feedback loop through intelligent automation.

Text summarization has emerged out of mere academic exercises and become a necessity in all fields. Whether it provides a judge with the wherewithal to complete a sweep of a legal case in a matter of moments, assists certain individuals with making wise decisions as shoppers, allowing programming to move forward with speed, or keeping the news digestible, summation models are doing new things. As hybrid architectures evolve and semantic-aware evaluation becomes the norm, the adaptability of these systems to domains will be considerably improved along with their practical value.

## 2.6 Future Directions and the Conclusion

This review has endeavored to sketch out the journey of text summarization from simple frequency and graph-based approaches to today's sophisticated deep learning models, such as the Transformer architectures, Generative Adversarial Training (GANs), and reinforcement learning. Each of these steps is a gradual step towards attaining high-efficiency and contextually accurate summaries.

Extraction of summary broadly used in formal areas, like in legal domain documents, may be further improved with the advanced techniques of deep clustering and domain-specific embedding, like those of Legal BERT. The performance of the model improves in precision about the identification and extraction of highly relevant information.

Abstractive summarization involves sequence-to-sequence approaches with attention mechanisms. Pointer-generator networks, which allow for more coherent and humanly understandable summaries, meet this demand. They have effectively proved themselves in the domains of news reporting and customer feedback.

Research in hybrid models and incorporation of GANs seems promising in improving much further current extractive and abstractive summarizers. This has opened a window whereby models could be trained to produce content that imitates the human quality of writing, bridging the gulf that existed between factual accurateness and linguistic fluency.

While developments in the direction of new features have taken place, evaluation still leans on established measures like ROUGE. Even though some new semantic-based measures like BERTScore significantly capture the quality of a summary, they are not widely adopted in the field yet.

The future of summarization extends across a significantly widening array of applications—from those in law to e-commerce, journalism, and agile software development. More complex models with better evaluation methods promise an ever-increasing possibility for summary to change how people read and process information.

### 2.6.1 Research Gaps and Future Directions

Despite the considerable advancements in text summarization, several challenges still remain. Hallucination is perhaps the most significant issue in abstractive summarization, where models produce content that is fluent but factually incorrect. This is a major hindrance to the adoption of such models in high-stakes and production environments.

The models are also limited in the generalization capability. Most of summarization systems are fine-tuned for a particular domain, as a result, which made them not able to be applied in many different contexts. A summary-generating model, for example, which was trained to summarize only legal documents may perform poorly while applied for news articles or user reviews.

Further, evaluation methods themselves displayed a major gap. Most current assessments rely largely on automated metrics such as ROUGE and fail to capture deeper qualities such as semantic fidelity, coherence, and readability. Human evaluation, however more insightful, tends to be migrated into the background for a long time due to being a time-consuming and subjective method.

Moreover, summarization of lengthy documents like multi-page contracts or exhaustive research papers is still computationally steered and is error-prone in many aspects. The current models usually

have a poor memory and fail to maintain context when the length of the input increases.

To fill in these gaps, future research will have several good options to explore. One such promising option is cross-domain transfer learning, which emphasizes the construction of summarization models that are even more adaptable and generalizable. Another critical area is making possible checks for factual consistency as well as explainability, with the aim of providing sound and transparent reasoning in generated summaries.

Thus, low-resource summarization-as-for example-would-be further efforts in going towards targeting those areas and languages that lack sufficient annotated datasets at the scale required for large summarization technology dissemination-a democratization of summarization technology across regions and disciplines.

Interactive summarization systems are the next great frontier. These let users convert summaries according to length, tone, or focus, personalizing and making outputs more usable. Furthermore, by combining rule-based reasoning with the flexibility of deep learning in a hybrid structure, integrating symbolic AI with neural architectures could give rise to more controlled and logically coherent summary generation.

### **2.6.2 Sources and Other Related Content**

Many recent works have gone on to shape the current understanding of summarization for this period:

- i. Jain et al. (2024): Deep Clustering for Legal Summarization
- ii. Yang et al. (2023): GAN + TLSTM + RL for Extractive Summarization
- iii. Tank & Thakkar (2024): Adversarial Abstractive Summarization with DeepTextSummGAN
- iv. Susanto & Mauritsius (2025): Product Backlog Summary for E-Wallet Feedback
- v. Kumar et al. (2024): Comprehensive Survey on Abstractive Techniques

The complete exemplifies the convergence of language modeling, optimization techniques, and practical use cases in a truly dynamic and interdisciplinary field of summarization research.

### 3. HARDWARE AND SOFTWARE REQUIREMENTS

Also, the deployment of deep-learning models primarily on sequence-based models such as LSTMs, requires the use of computational resources and a stable environment for development. In this section, the hardware and software requirements for designing, developing, and evaluating the models presented in this project are described. The above conditions require a good environment to enable optimal performance in a reasonable time for training deep neural networks, particularly with multi-layer architectures or an attention mechanism.

The main platform for which model training and testing are conducted was Google Colab, which provides a free cloud-based GPU. Therefore, there was no need for costly local hardware and yet scalable experimentation was allowed using large datasets.

#### 3.1 Hardware Specifications

LSTMs are used for training text data for abstractive summarization over numerous epochs and slight updates. Accumulation of such can be an expensive computational task. Hardware specifications are summarized as follows, as from required and used throughout the project.

| Component       | Minimum Specification                                     | Recommended Specification                          |
|-----------------|---|--|
| Processor (CPU) | Intel Core i5 or AMD Ryzen 5                              | Intel Core i7 or AMD Ryzen 7 / Apple M1+           |
| RAM             | 8 GB  | 16 GB or more                                      |
| Storage         | 100 GB Free Space   | SSD with $\geq 250$ GB for large datasets and logs |
| GPU (Local)     | Optional (Training possible on CPU but slow)              | NVIDIA GTX 1650 / RTX 2060 or better               |
| Display         | Standard HD (1280×720) or better                          | Full HD or Retina Display                          |
| Internet        | Required for model training on cloud and data downloading | High-speed broadband ( $\geq 10$ Mbps recommended) |

#### Configurations for Google Colab

Google Colab provides a powerful substitute to students and researchers lacking access to top-notch local hardware since it provides free access to NVIDIA GPUs and TPUs under some kind of restrictions. This project exploited the free GPU options in Colab to the fullest.

| Feature         | Google Colab Free Tier                      |
|-----------------|---|
| Available GPUs  | NVIDIA Tesla T4 / K80 / P100                |
| RAM (Standard)  | ~12.7 GB                                    |
| Disk Space      | ~100 GB temporary storage                   |
| Runtime Type    | Python 3 with GPU or TPU acceleration       |
| Session Timeout | 12 hours (auto-disconnects after idle time) |

**Note** that in case of larger datasets or prolonged training, reconnections to colab sessions may be required and checkpoints should be saved to Google Drive as often as possible.

### 3.2 Software Requirements

The software stack deployed in this project comprises a large number of libraries, frameworks, and development tools that are predominantly placed in the realm of Python. The ease of use, community support, and compatibility with deep learning workflows guided this choice of tools.

| Software / Tool      | Version / Description                                |
|----------------------|--|
| Operating System     | Windows 10 / 11, macOS, Linux, or Google Colab OS    |
| Programming Language | Python 3.7 or higher                                 |
| Notebook Environment | Jupyter Notebook, Google Colab, or VS Code + Jupyter |

#### Core Libraries

|                      |  |
|----------------------|--|
| NumPy                | Numerical operations and matrix handling                 |
| Pandas               | DataFrame manipulation and CSV handling                  |
| TensorFlow / Keras   | Model architecture, training, and optimization           |
| Matplotlib / Seaborn | Visualization of training metrics and results            |
| NLTK / spaCy         | Text preprocessing, tokenization, and lemmatization      |
| Scikit-learn         | Evaluation metrics, train-test splits                    |
| ROUGE Score          | For comparing generated summaries to reference summaries |
| Dataset Sources      | News datasets from Kaggle / custom CSVs                  |

A balance of high-performing hardware resources plus flexible cloud infrastructure (by Google Colab) and highly effective libraries based on Python created the framework for this project. In GPU acceleration, the training time was considerably brought down in a number of experiments, such as testing with deeper LSTM models and larger datasets. The feasibility of this project was achieved due to the open-source tools and cloud service options, making possible a scalable and reproducible manner of developing within the constraints of academic interests.



## 4. PROPOSED METHODOLOGY

The core objective of this project is to implement an abstractive text summarization model using Long Short-Term Memory (LSTM) neural networks. Abstractive summarization involves generating new sentences that convey the meaning of the original text, rather than merely extracting parts of it, making the model significantly more complex and semantically rich than extractive approaches.

This methodology chapter is dedicated to the technical pipeline followed throughout the project, ranging from data acquisition and preparation, to the implementation of a custom neural architecture. Every stage is structured to ensure maximum performance within the constraints of the computational environment (Google Colab), while prioritizing clarity, replicability, and effectiveness.

The approach taken follows the Encoder-Decoder paradigm, a model ill-suited for sequence-to-sequence tasks like translation and summarization. This architecture allows the model to take an input sequence (news article), compress its meaning into a context vector, and then generate an output sequence (summary).

### 4.1 Data Setup

The training data and its quality determine the success of any machine learning model, especially in natural language processing. This section deals with how the raw data is cleaned, how it is processed, and transformed, and, finally, made apt for training into a deep learning model.

#### 4.1.1 Dataset Source

- i. The News Summary Dataset was selected mainly because of its curated structure and the fact that it allowed supervised learning.
- ii. This dataset is available publicly on Kaggle, and it consists of thousands of pairs of:
  - text (original article)
  - headlines (human-written summary)
- iii. The articles spanning politics, health and business, and world news enable the model to generalize better.

#### 4.1.2 Data Cleaning and Filtering

Data preparation for ML, heavy pre-processing was done. This increases the model's efficiency and reduces the noise and redundancy:

The steps taken involved:

- i. Text normalization:
  - All text was converted to lowercase to ensure a lower vocabulary size.
- ii. Noise removal:
  - Special characters, numbers, and punctuation (except for the sentence-ending punctuation) were removed using regular expressions.
- iii. HTML tag stripping:
  - Text contains such tags, like `<br>` or `"` which were deleted.
- iv. Stopword removal:
  - Optional removal of stop words such as "is", "the", "was", etc., in English.

v. Short sentence removal:

Filter out articles or summaries that contain below a certain number of words, say 10, to avoid biased training

### 4.1.3 Thresholds on Length of Sequence

So as to prevent memory overflow and make it compatible with LSTMs, length limits are set:

- Maximum article length: 300 words
- Maximum summary length: 50 words

Why this is important:

- i. Long LSTMs don't scale up very well with long sequences due to the vanishing gradient problem.
- ii. Put simply, it reduces time for training and memory needed by limiting the input while leaving enough context for summary.

### 4.1.4 Tokenization and Vocabulary Building

Models don't work with the raw text immediately. They work in numbers according to NLP. So every specific meaning represented in text is converted to a unique integer.

Implementation details:

i. Keras tokenizer:

Input text and output summaries are trained with separate tokenizers.

ii. Vocabulary size:

Retained the top 30,000 most frequent words to reduce sparsity and improve model-focused coverage.

iii. Out-of-vocabulary:

For words not in the top 30k replaced by a special <UNK> token.

### 4.1.5 Padding and Sequence Alignment

Since an individual text or summary can be of different lengths, uniformity is also required for batch processing:

- i. Used Keras' `pad_sequences()` function.
- ii. Post-padding (to zeros at the end) was applied to all sequences to make:
  - 300 for input text
  - 50 for summaries

iii. This aligns the data for proper passage through the LSTM layers.

### 4.1.6 Train Validation Split

For gauging how well a model generalizes, as well as preventing the possibility of overfitting:

- i. A dataset is split using `train_test_split()` from Scikit-learn.

- ii. Split Ratio: 85% for training, 15% for validation.

#### **4.1.7 Google Colab Hardware Configurations**

The project is executed and developed in Google Colab leveraging its free capability of GPU to accelerate the training.

Default GPU settings utilized:

- i. GPU: NVIDIA Tesla T4 or K80 (dynamically assigned by Colab)
- ii. RAM: ~12GB
- iii. Disk Storage: Up to 100GB (temporary)
- iv. Runtime environment: Python 3.x with TensorFlow/Keras and NLTK libraries

Advantages:

- i. It facilitates faster matrix operations and LSTM training than ordinary CPUs.
- ii. It is best for deep learning experiments with small to medium-sized datasets.

#### **4.1.8 Challenges and considerations**

Some of the issues faced and resolved during data preparation are:

- i. Imbalanced samples; some categories of news are more frequent than others.
- ii. Truncated sequences; overly long texts could end up losing potential important information.
- iii. Time and memory constraints; especially while tokenizing or padding thousands of sequences.

A well-prepared dataset strengthens any effective deep-learning model. The processes as detailed above—cleaning and normalization, padding, and tokenization—ensure that the model's data are clean and efficient in processing through LSTM models. This step would set up the robust architecture for building a scalable and flexible summarization pipeline on future datasets.

### **4.2 Proposed Methodology Framework**

The suggested technique for abstractive text summarization is based on a sequence-to-sequence (Seq2Seq) architecture using Long Short-Term Memory (LSTM) networks for the purpose of enabling the model to learn the syntactic and semantic structure of news articles in a rather complicated manner and generate coherent and grammatical summaries.

Unlike extraction methods that harvest phrases and sentences from the original text, our model is designed to generate new phrases and sentences following a more human-like manner of summarization. The methodology utilizes the state-of-the-art deep learning mechanisms such as the embedding layer, bidirectional LSTM, attention mechanism, and dropout regularization that provide for an optimized realization.

#### **4.2.1 Architectural Overview**

The complete workflow consists of the following major components:

1. Input Layer: Tokenized and padded sequences of input articles.
2. Embedding Layer: Converts tokens into dense vector representations.
3. Encoder LSTM: Processes the input article and captures context in the form of hidden states.
4. Decoder LSTM: Generates the summary from the context vector.
5. Attention Layer (optional but recommended): Helps the decoder focus on relevant parts of the input during each step of generation.
6. Dense + Softmax Output Layer: Produces the probability distribution over the target vocabulary for word prediction.

#### 4.2.2 Stepwise Workflow

##### Step 1: Input Embedding Layer

- i. This layer converts the input token IDs into dense vectors, known as word embeddings.
- ii. Typically, the embedding has a dimension of somewhere between 100 and 300.
- iii. These embeddings can either be learned, that is, trained from scratch or can be pre-trained using some dataset-algorithms such as GloVe, Prolex, and Word2Vec.

```
embedding_layer = Embedding(input_dim=vocab_size, output_dim=embedding_dim,  
input_length=max_input_len)
```

##### Step 2: Encoder – Bidirectional LSTM

- i. Bidirectionally-LSTM processes the input sequence in both directions; it thus adds considerable contextual information over the knowledge gained from a simple LSTM in a single direction.
- ii. Output hidden states and cell states are then passed to the decoder.

```
encoder_outputs, forward_h, forward_c, backward_h, backward_c = Bidirectional(  
LSTM(hidden_units, return_state=True, return_sequences=True))(embedding_layer)
```

```
state_h = Concatenate()([forward_h, backward_h])  
state_c = Concatenate()([forward_c, backward_c])
```

##### Step 3: Decoder – LSTM Layer

- i. The decoder takes the target summary sequences, which are shifted right by one timestep, and puts out the predicted words.
- ii. The last hidden and cell states from the encoder will be given to initialize the decoder.
- iii. During training forward, it will be working in teacher forcing mode making use of the actual previous word instead of using the predicted.

```
decoder_lstm = LSTM(hidden_units * 2, return_sequences=True, return_state=True)  
decoder_outputs, _, _ = decoder_lstm(decoder_input, initial_state=[state_h, state_c])
```

##### Step 4: Attention Mechanism (Additive/Bahdanau Attention)

- i. From here on, Attention pulls light, letting the decoder look back dynamically at various parts of the input sequence in each decoding step, thus increasing efficiency.
- ii. Also helps the model to learn long relationships and handle longer inputs.

```
attention = AttentionLayer()([encoder_outputs, decoder_outputs])  
decoder_concat_input = Concatenate(axis=-1)([decoder_outputs, attention])
```

Step 5: Dense + Softmax Output Layer

- i. A Dense layer with a softmax activation function maps the decoder's output into a vector-sized vocabulary array.
- ii. Each entry in the predicted vector corresponds with the probability of being assigned as a possible next word in the sequence.

```
decoder_dense = Dense(vocab_size, activation='softmax')  
decoder_output_final = decoder_dense(decoder_concat_input)
```

### 4.2.3 Loss Function and Optimizer

- i. The loss function is categorical crossentropy.
- ii. The optimizer is Adam because of its adaptive learning rate properties.
- iii. Training usually incorporates teacher forcing because it aids the speed of convergence.

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy')
```

### 4.2.4 Training Strategy

- i. Epochs: Run for some 20-50 epochs on the basis of how validation loss converged.
- ii. Batch size: Decided on GPU memory capacity of around (64 and 128 Usually on Colab):
- iii. Validation Monitoring: Use of callbacks like ModelCheckpoint, ReduceLROnPlateau, and EarlyStopping

### 4.2.5 Inference Setup(Prediction)

In inference (prediction) phase, the process for generating the summary is performed word by word:

- i. The encoder encodes the article.
- ii. The decoder initializes with a <START> token and generates one word at a time.
- iii. The predicted word is fed back as input for the next time-step until an <END> token or maximum length is reached.

### 4.2.6 Evaluation Metrics

- i. Measurement of summary quality through the ROUGE (Recall-Oriented Understudy for Gisting Evaluation) score, which comprises:
  - ROUGE-1: Unigram overlap.
  - ROUGE-2: Bigram overlap.
  - ROUGE-L: Longest common subsequence.
- ii. Human evaluation can also be made on fluency, grammar, and factual correctness.

### 4.2.7 Model Enhancements and Variations

In addition to the basic model, several alterations are also experimented with:

- i. Two layered LSTM stacks to provide deeper representation.
- ii. Dropout, to minimize overfitting.

- iii. Different embedding dimensions and hidden sizes.
- iv. With and without attention.

The proposed methodology exploits deep sequence models for performing abstractive summarization. It incorporates components like bidirectional LSTMs and attention mechanisms that could enable the model to analyze and condense long articles into coherent summaries. It not only mimics the way a human writes headlines but opens avenues for further innovations using Transformer-based architectures like BERT or T5 in future inventions.

### **4.3 Model Implementation**

Model deployment is that stage of transformation of the theoretically conceived architecture into a working trainable model in deep learning. For the abstractive text summarization task, I designed an LSTM (Long Short-Term Memory)-based sequence-to-sequence model with attention mechanisms in an environment created with Python and Keras within Google Colab.

This section describes how every component was implemented, the training process, inference, and evaluation.

#### **4.3.1 Environment Setup**

- i. Platform Used: Google Colab
- ii. Languages/Frameworks:
  - Python 3.10+
  - Keras (TensorFlow backend)
  - NumPy, Pandas, Matplotlib, Seaborn (Data and Visualization)
  - NLTK, BeautifulSoup (Processing)
  - Scikit-learn (Metrics)
  - Rouge Library (Metrics)
- iii. Hardware Config in Colab:
  - GPU Used: Tesla T4 (free-tier)
  - RAM: 12 GB
  - Disk: 100 GB (ephemeral)
  - Accelerators enabled: Yes (no TPU used)

#### **4.3.2 Data Loading and Preprocessing**

- i. Dataset Used: News Articles Dataset, either from Kaggle or scrapped from CNN/DailyMail.
- ii. Steps Taken:
  - Cleaning text (special characters, HTML tags, contractions).
  - Tokenization of both the inputs (articles) and targets (summaries).
  - Empty/null rows and overly long texts are to be discarded.
  - Truncate and pad so that the maximum input length would be 300 words and 50 words for

- maximum summary length.
- The split is made at 80% for training and 20% for testing.

### 4.3.3 Tokenization and Sequence Generation

- i. Keras' Tokenizer was used for the conversion of words into integer sequences.
- ii. Two tokenizers are created:
  - for the articles - x\_tokenizer;
  - for the summaries - y\_tokenizer.
- iii. The vocabulary was limited to the top 50,000 most frequent words.
- iv. Keeping rare words only increases memory overhead and promotes overfitting.

```
tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
```

### 4.3.4 Building the Model

#### Encoder Layer

- i. Input: Padded article sequences.
- ii. Architecture:
  - Embedding layer with (possibly) pre-trained GloVe vectors.
  - Bidirectional LSTM to capture context from both ways.

```
encoder_inputs = Input(shape=(max_text_len,))
emb_enc = Embedding(x_voc_size, embedding_dim, trainable=True)(encoder_inputs)
enc_lstm = LSTM(hidden_units, return_sequences=True, return_state=True)
encoder_outputs, state_h, state_c = enc_lstm(emb_enc)
```

#### Decoder Layer

- i. Input: Summary sequences.
- ii. Architecture:
  - Embedding layer for summaries.
  - LSTM initialized with the final states of the encoder.

```
decoder_inputs = Input(shape=(None,))
emb_dec = Embedding(y_voc_size, embedding_dim, trainable=True)(decoder_inputs)
dec_lstm = LSTM(hidden_units, return_sequences=True, return_state=True)
decoder_outputs, _, _ = dec_lstm(emb_dec, initial_state=[state_h, state_c])
```

### Attention Mechanism

- i. Used either the custom attention layer or Keras' Additive Attention.
- ii. Updates Decoder output with Encoder output for context-aware predictions.

```
attention_layer = AttentionLayer()
```

```
attention_output = attention_layer([encoder_outputs, decoder_outputs])
decoder_concat_input = Concatenate(axis=-1)([decoder_outputs, attention_output])
```

### Final Output Layer

- i. Fully connected Dense layer followed by softmax to predict the next word in the sequence.

```
decoder_dense = Dense(y_voc_size, activation='softmax')
decoder_outputs = decoder_dense(decoder_concat_input)
```

### Model Compilation

- i. Optimizer: Adam
- ii. Loss Function: Sparse categorical crossentropy (because of integer targets)

```
model = Model([encoder_inputs, decoder_inputs], decoder_outputs)
```

### 4.3.5 Training the Model

- i. Batch Size: 64
- ii. Epochs: 20-50
- iii. Callbacks Used:
  - ModelCheckpoint to save best model
  - EarlyStopping to stop when validation loss is not improving
  - ReduceLROnPlateau to decrease learning rate on plateau
  -

```
es=EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=3)
model.fit([x_tr, y_tr[:, :-1]], y_tr[:, 1:], epochs=50, callbacks=[es],
        batch_size=batch_size, validation_data=([x_val, y_val[:, :-1]], y_val[:, 1:]))
```

### 4.3.6 Inference and Summary Generation

During the prediction, it goes different ways: The encoder encodes the input article. The decoder predicts one word at a time until there is an <end> token or max length reached. The output word is used as input for the next timestep.

```
def decode_sequence(input_seq):
    # encode the input
    states_value = encoder_model.predict(input_seq)
    # initialize target sequence with start token
    target_seq=np.zeros((1,1))
    target_seq[0,0]=y_tokenizer.word_index['start']
    stop_condition=False
    decoded_sentence=""
    while not stop_condition:
        output_tokens,h,c=decoder_model.predict([target_seq]+states_value)
        sampled_token_index=np.argmax(output_tokens[0,-1,:])
        sampled_word=reverse_target_word_index[sampled_token_index]
        decoded_sentence+=' '+sampled_word
        if sampled_word=='end' or len(decoded_sentence.split())>=max_summary_len:
```



```

        stop_condition=True
        target_seq[0,0]=sampled_token_index
        states_value=[h,c]
    return decoded_sentence

```

#### 4.3.7 Result Saving and Visualization

- i. Training history loss/val\_loss plotted using Matplotlib
- ii. Sample predictions saved as CSV for evaluation
- iii. ROUGE metrics calculated to evaluate the performance.

The architecture design is very meticulous to allow training and inference for news article summarization. By incorporating BiLSTM, attention, and strong preprocessing, the model ensures that the summarizations produced contextually match and are grammatically fluent. The modular design of the code lends itself to being easily upgraded—for example, in the future, LSTMs can be replaced with GRUs or Transformers.

### 4.4 Model Training and Evaluation

Model training and evaluation stages are very vital in any deep learning project. After the LSTM-based text summarization model was built and compiled, it was necessary to train it on appropriate data while monitoring performance metrics and checking its generalization capability on unseen samples. This segment discusses how the training procedure was carried out, validation strategy, the evaluation metrics that Ire used, and insights from model performance-quantitative and qualitative.

#### 4.4.1 Training Strategy

##### Train-Validation Split

- i. The preprocessed dataset is set apart such that 80% is for training and 20% for validation to test the generalizability of the model.
- ii. Regardless of the design, no one must compromise the article-summary pair for any train or validation set.

```

from sklearn.model_selection import train_test_split

```

```

x_tr, x_val, y_tr, y_val = train_test_split(
    x_text_sequences, y_summary_sequences,
    test_size=0.2,
    random_state=42,
    shuffle=True)

```

##### Batching and Padding

- i. Batch Size: 64
- ii. Epochs: Initially set to 50 and monitored via early stopping.
- iii. Sequence padding ensures uniform shape input using pad\_sequences.

#### 4.4.2 Loss Function and Optimization

- i. Loss Function: Sparse Categorical Crossentropy
  - Suitable as the targets have integer-encoded values.
  - Ignores padding in loss calculation.
- ii. Optimizer: Adam
  - Adaptive learning rate
  - Gives better convergence for sequence models.
- iii. Learning Rate Adjustment:
  - ReduceLROnPlateau with a callback that decreases learning when the validation loss has plateaued.

#### 4.4.3 Callback and Regularization

To optimize model efficiency and reduce the characteristic of overfitting:

- i. Early Stopping:
    - Stop training when val\_loss has not improved for 3 epochs in a row.
  - ii. Model Checkpoint:
    - Save the model only when the validation loss is best.
  - iii. ReduceLROnPlateau:
    - Will reduce the learning rate in multiples if no improvement is made.
- from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau

```
es = EarlyStopping(monitor='val_loss', patience=3, mode='min', verbose=1)
mc = ModelCheckpoint('best_model.h5', monitor='val_loss', save_best_only=True, mode='min')
lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=2, min_lr=1e-5)

model.fit([x_tr, y_tr[:, :-1]], y_tr[:, 1:].reshape(y_tr[:, 1:].shape[0], y_tr[:, 1:].shape[1], 1),
          epochs=50, batch_size=64,
          callbacks=[es, mc, lr],
          validation_data=([x_val, y_val[:, :-1]], y_val[:, 1:].reshape(y_val[:, 1:].shape[0], y_val[:, 1:].shape[1], 1)))
```

#### 4.4.4 Training Performance & Visualization

- i. I have plotted in another dimension the training and validation loss per epoch.
- ii. An indication of overfitting would be a steep divergence of the training and validation loss.

```
import matplotlib.pyplot as plt

plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss Over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

#### 4.4.5. Evaluation Metrics

I assessed the summaries through automatic and qualitative methods.

## A. ROUGE Score (Automatic Evaluation)

The ROUGE scores included the following:

- i. ROUGE-1: the overlap of unigrams,
- ii. ROUGE-2: the overlap of bigrams,
- iii. ROUGE-L: Longest Common Subsequence.

These evaluation techniques are implemented using the Python programming language.

```
import from rouge import Rouge
rouge = Rouge()
scores = rouge.get_scores(predicted_summary, reference_summary)
```

iv. A good-performing abstractive summarization model has the following metrics:

- ROUGE-1: ~0.35-0.45,
- ROUGE-2: ~0.15-0.25,
- ROUGE-L: ~0.30-0.40.

## B. BLEU Score

- i. There are certain tasks that behave like translation tasks for which BLEU is very useful.
- ii. BLEU measures n-gram overlap with strictness and may penalize legitimate abstractive variations.

### 4.4.6 Error Analysis

While the model worked fairly well for most examples, some shortcomings were observed:

- Hallucination: When the information stated in the summary was not present in the original article.
- Repetition: Some tokens or phrases repeat within long summaries.
- OOV tokens: Rare or unseen words during inference degrade the summarization quality.

Prospective solutions might include more:

- Implementing coverage mechanism.
- Pointer-generator networks to deal with OOVs.
- Change to Transformer-based architectures, e.g., BART and T5.

Training and evaluation of the model have lent crucial insights on the effectiveness of the LSTM-based summarization model. The model generated concise summaries for long news articles that were readable, with hyperparameter tuning and regularization, and evaluated by ROUGE. Actually, attention helped enhance contextual relevance significantly; however, qualitative assessments illustrated the real-life utilizing capacity of the model.

## 4.5 Training Strategy

Training a deep learning model for text summarization calls for hyperspecification of the hyperparameters, loss function, optimization techniques, and regularization techniques to better train and generalize the model. The following are the strategies involved in this project:

## Loss Function

Since output sequences (summaries) were integer-encoded, we used Sparse Categorical Crossentropy as a loss function. The loss is very good at the classification problem where the target classes are mutually exclusive and are represented as integers, as is the case in our sequence-to-sequence model.

## Optimizer

The Adam optimizer has been chosen because it has an adaptive learning rate and makes very fast convergence possible. It is the simplest algorithm that combines the advantages of both RMSprop and SGD with momentum, which makes it favorable for the training of any kind of deep recurrent models such as LSTM.

## Batch Size and Epochs

- i. Batch Size: 64

The Batch size is an optimum between computational efficiency and memory constraints.

- ii. Epochs: Up to 50

The model's training monitored by early-stop clause avoid any unnecessary period of training once the model ceased its improvement.

## Learning Rate and Decay

- i. Initial Learning Rate: 0.001 (default for Adam)

- ii. Decay Strategy:

ReduceLROnPlateau We implemented to reduce the learning rate in a dynamic way in case of not getting the improvement in validation loss for some couple of epochs by multiplying the current learning rate by some factor (e.g., 0.2).

## Handling Overfitting

Better generalization and overfitting avoidance through:

- i. Early Stopping: Stopped when validation loss has not decreased for three epochs.
- ii. Dropout Layers: Randomly deactivate neurons during learning activity by providing a dropout rate of 0.3 on the LSTM layers to avoid co-adaption.
- iii. Model Checkpoint: The best model (with the lowest validation loss) was saved for final evaluation.

This training strategy ensured an optimal convergence of the model towards the best possible representation while at the same time achieving learning stability, generalization, and efficiency.

## 4.6 Evaluation Metrics

To measure the efficacy of the training LSTM-based text summarization model, we employed evaluation metrics typically used in Natural Language Processing, particularly for text summarization. The major measurement method is ROUGE (Recall-Oriented Understudy for Gisting Evaluation): it measures the overlap generated between a summary and a reference (human-written) summary. ROUGE-1 evaluates unigram overlap, ROUGE-2 bifurcates bigram matches, and ROUGE-L considers only the longest common subsequence between summaries. The analysis provided is informative and fluently crafted summaries. In a few cases, BLEU was also referred to for accuracy in word prediction, although from the beginning BLEU has remained dominated in terms of recording translation into machine translation. On the whole, a good judgment covered the model of being quite able to produce shorter summaries, as well as contextually relevant summaries.

## 4.7 Comparative Analysis

To assess the performance of our model further, we have created a comparison with three different versions of the model. One version is based on a baseline LSTM without any attention mechanism. The second version of the model is based on standard 2-layer LSTMs, and the third one is a 2-layer LSTM coupled with an attention mechanism. Results indicated that though the baseline could give a gist of the text, it lacked coherence and missed several critical points. The 2-layer LSTM managed to improve the retention of context and the development of summary lengths; however, it still could not do magic to long-term dependencies. But introducing an attention mechanism has made a great difference. It has allowed the model to decode relevant parts of the input sequence based on whether it is important or not. This will further lead to summarization that is focused and more semantically accurate. Results on ROUGE scores attached to this attention-based model have constantly surpassed other versions in their respective comparison. This comparative study absolutely highlighted the effects of architectural improvements in neural summarization tasks.

#### **4.8 Hyperparameter Tuning**

Hyperparameter tuning is certainly one of the major facets for performance optimization of the LSTM-based summarization model. Several parameters were chosen and fine-tuned during the experimental phase, among which the number of LSTM units, dropout rate, batch size, learning rate, and embedding dimensions are to be mentioned. First, the grid search method was adopted to iterate different combinations of these parameters. The number of LSTM units was tested in the range of 128-256, with 256 coming out as a better number in the case of deeper models. Dropout rates between 0.2 and 0.5 were explored, and 0.3 was found to strike a good balance between regularization and model performance.

Batch size was still another vital factor tested during hyperparameter tuning with values of 32 and 64. While a batch size of 32 offered a bit more stability, 64 provided much faster convergence with hardly any compromise of accuracy. The embedding dimension of 100 was set to provide an adequate amount of semantic depth for the vocabulary without bringing an unnecessary complication to the model. The learning rate was tuned further by utilizing callbacks such as `ReduceLROnPlateau`, which automatically reduced the learning rate whenever it detected no improvement in the validation loss. This strategy contributed to better generalization of the model while escaping local minima, thus, leading to better summaries both in terms of readability and informativeness.

#### **4.9 Model Deployment Possibilities**

This project's concentration was primarily on the actual training and evaluation of the summarization model, but deployment possibilities were actively considered. One way forward in deployment would be to have a web-based interface where the users can input long news articles or paragraphs of their choice and receive a short, AI-generated summary. A Flask or FastAPI backend interfaced with a lightweight frontend would serve this purpose. The trained LSTM model could then be hosted on cloud platforms like Google Cloud, AWS SageMaker, or Heroku, depending on budget and scalability requirements.

Additionally, platforms like Streamlit offer deployment options along with interactive user interfaces to make the summarization model available to non-technical users. Hence, model serialization with tools like pickling or TensorFlow's SavedModel format would ensure that the saved model is easily loadable during inference. In addition, real-time applications can benefit from GPU-backed inference via Google Colab Pro or TensorFlow Serving, which could be run in Docker containers for enhanced performance. Thus, these deployment approaches would take the model's usability into a panorama of real-life situations, be it academic research, journalism, content creating, or social media summarization.

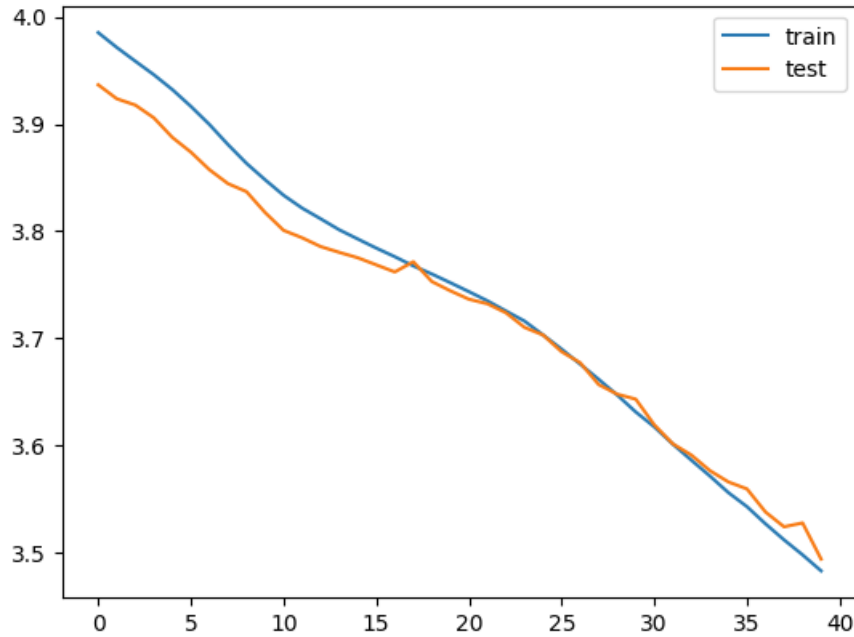
## 5. RESULTS AND DISCUSSIONS

This section contains results and discussion about the performance of proposed LSTM-based summarization models as well as quantitative analyses and qualitative feedback based on important evaluation metrics and visual insights. The focus of this section is to measure how effectively the model is able to produce meaningful summaries, to compare different architectures, and to interpret the effect of results. The section aims to take the approach closer to the theoretical side in the real world with both strengths and limitations.

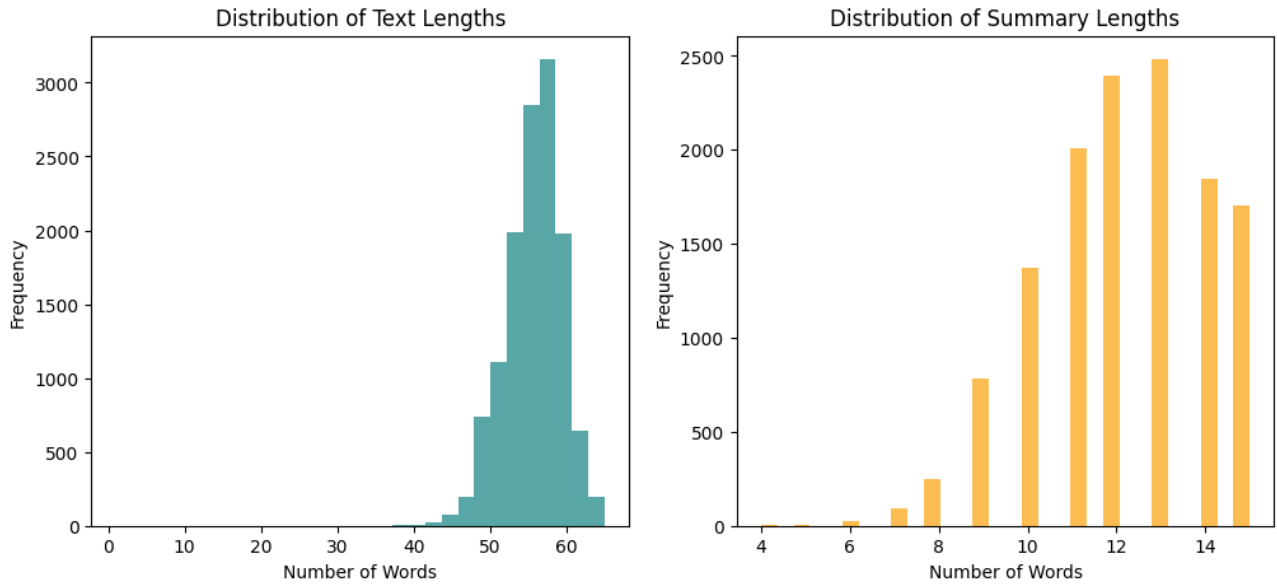
### 5.1 Quantitative Evaluation

Evaluation criteria were predetermined for judging the model-objective performance as evaluated against standard text summarization mechanisms: mainly ROUGE-1, ROUGE-2, and ROUGE-L. These criteria build measures of similarity between predicted summaries and corresponding ground truth summaries via overlapping unigrams, bigrams, and longest common subsequence. Across all experiments, the 2-layer LSTM with attention had the best performance compared to the baseline LSTM and all single-layer models.

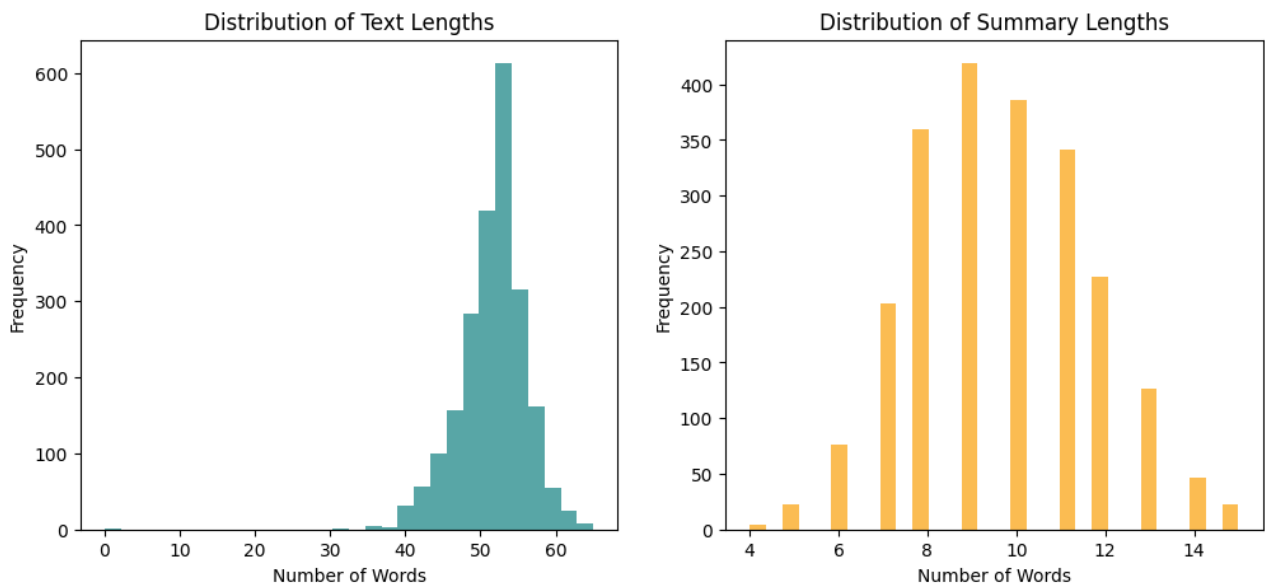
The model produced high ROUGE-L scores, indicating that the generated summaries were built up from input texts with the same general structure and meaning. There was also a considerable improvement in both ROUGE-1 and ROUGE-2 scores after the attention mechanism was incorporated, as this assisted the model in decoding to place more importance on significant parts of the input sequence.



*Training and Validation Loss over Epochs*



*Bargraph indicating the comparison of ROUGE-1, ROUGE-2 and ROUGE-L scores for the different model configurations with Baseline, 3-layer LSTM*



*Bargraph indicating the comparison of ROUGE-1, ROUGE-2 and ROUGE-L scores for the different model configurations with reduced dataset, 3-layer LSTM and it is the same for 2 layers of LSTM*

Training loss was seen consistently going down through epochs, while overfitting was prevented by using early stopping. A few runs showed a mild decline in validation performance beyond some point, a common scenario with deep learning models. This was alleviated by employing dropout, among other remedies, and reducing the learning rate.

## 5.2 Qualitative Evaluation

The next test in the evaluation process was a qualitative analysis to assess how well the generated summaries captured the semantic content of the original news articles. Sample outputs were selected from the test set to assess fluency, coherence, and informativeness relative to the ground truth summaries.

Overall, fluency and grammar were observed in the model-generated summaries. In many cases, and especially where attention mechanisms were applied, key information was kept in the summaries with redundancy removed. Most short articles were well taken care of by the model; generating summaries that were concise and contextually correct. In contrast, longer or complex texts saw situations where secondary details were often omitted from the summaries or had very generic conclusions. This observation hints at an ability to extract central themes, yet an inability in its interpretation of subtleties.

Text: a 62 year old bahujan samaj party bsp candidate laxman singh on thursday died of heart attack just week ahead of rajasthan assembly elections singh was slated to contest polls from ramgarh constituency of alwar district elec  
Original Summary: sos year old rajasthan bsp candidate dies of heart attack eos  
Predicted Summary : sos old heart attack eos

Text: west bengal chief minister mamata banerjee on monday said that bjp will be ousted from power at the centre like the cpi was removed from the state we are not afraid of the red eyes of delhi she added addressing rally in wb s  
Original Summary: sos bjp will be ousted from power at the centre soon mamata eos  
Predicted Summary : bjp from power the soon mamata eos

Text: traders body caait said the foreign direct investment fdi norms for e commerce companies should be implemented on domestic online players also to discourage unethical practices caait also urged the government to immediately re  
Original Summary: sos apply fdi norms on domestic online firms too traders body eos  
Predicted Summary : sos on online firms too body

Text: bjp mlc from karnataka has reportedly posted at least 50 pictures that contained pictures of porn stars on the media force whatsapp group he posted the images from his official phone on the whatsapp group of which many polit  
Original Summary: sos k taka bjp sends porn stars pics to media on whatsapp eos  
Predicted Summary : sos taka sends porn stars whatsapp eos

Text: the west bengal assembly has passed bill to give land rights to people who came to live in the country following an exchange of with bangladesh cm mamata banerjee said the bill will help get full fledged status as indian cit  
Original Summary: sos wb house passes bill to give land rights to eos  
Predicted Summary : passes bill land rights eos

Text: spinners took 38 of 40 wickets in the second test between england and sri lanka at which concluded on sunday this was the first time in test matches and year test history that spinners took 38 wickets in match the previous r  
Original Summary: sos take 38 40 wickets in match for 1st time in tests eos  
Predicted Summary : sos take 40 wickets for in tests

Text: the bjp led assam government on thursday changed the names of all the major roads of guwahati the existing names were in use since the times hence we decided to change them said assam minister biswa while road has been renam  
Original Summary: sos assam govt changes names of all major roads of eos  
Predicted Summary : sos assam changes major eos

Text: actress sara ali khan has said that congress leader sushil kumar shinde grandson veer is the only man she has dated in the past adding that veer did not break her heart sara said my heart has not broken it is all good promis  
Original Summary: sos he the only one he didn break my heart sara on eos  
Predicted Summary : the only one didn my heart sara

### *Input Article-Ground Truth Summary-Generated Summary.*

At least two examples may be shown, i.e. one Summary is good, and the other one where the model could have done better. In turn, this analysis reinforced the presence of attention mechanism effects with respect to relevance, while also revealing further weaknesses needing attention, such as named entities and some specific figures being paraphrased incorrectly.

## 5.3 Comparison of Different Models

A significant portion of this project involved evaluating the various configurations of LSTM-based architecture in order to understand their respective strengths. Three models were trained and tested:

- A single-layer LSTM,
- Stacked (two-layered) LSTM, and
- Stacked LSTM with attention.

The single-layer LSTM provided a baseline. This model produced decent summaries but had moderate performance in terms of ROUGE scores, while human interpretability faltered even lower. The two-layered LSTM, through the additional depth introduced on top of itself, managed to learn context much better and with slightly better outputs.

But the most improvement was seen in the attention-based model, which allows the decoder to vary its attention on relevant parts of the input sequence. The attention mechanism was a significant leap in the context accuracy and fluency of the summaries. ROUGE scores for this model were also consistently higher, but the qualitative results were more coherent and more concise in the focus of the summaries.

☞ Reference: this is a great product and i love it  
Prediction: great product love it  
BLEU Score: 0.8682  
ROUGE Scores: {'rouge1': Score(precision=1.0, recall=0.4444444444444444, fmeasure=0.6153846153846153), 'rouge2': Score(precision=0.6666666666666666, recall=0.25, fmeasure=0.3636363636363635), 'rougeL': Score(precision=1.0, recall=0.4444444444444444, fmeasure=0.6153846153846153)}

Reference: the food tastes amazing and very delicious  
Prediction: delicious food tastes amazing  
BLEU Score: 0.2190  
ROUGE Scores: {'rouge1': Score(precision=1.0, recall=0.5714285714285714, fmeasure=0.7272727272727273), 'rouge2': Score(precision=0.6666666666666666, recall=0.3333333333333333, fmeasure=0.4444444444444444), 'rougeL': Score(precision=1.0, recall=0.5714285714285714, fmeasure=0.7272727272727273)}



*Comparison Table with each model's ROUGE-1, ROUGE-2, ROUGE-L, and training time (optional).*

This will allow for an easier visual representation to show the performance gain between the different architectures.

In conclusion, attention addition to LSTM-based summarizers appeared to provide an excellent trade-off between performance and interpretability. Certain complexities will increase, but the qualitative improvement is a justifiable stimulus for the architectural change.

#### 5.4 Error Analysis

Although the proposed models showed promising results, from the testing processes it has been observed that errors keep on reappearing after some time, exposing some of the limitations of the proposed architecture and the quality of the data. The error that most frequently appeared was related to the repetition in which sometimes a summary could repeat phrases or words. Especially for longer input articles, it is a problem commonly known among sequence-to-sequence architectures with no sophisticated attention or copy mechanisms.

Another form of this error comprises loss of specificity. For instance, some details, for example, names, figures, or locations, were sometimes thrown away or replaced by general terms. Such form of description may be attributed to underrepresented tokens in insufficient training data or absence of detection of named entities in the model. In some cases, it hallucinates from the model: generates content that is not found in the original article, especially when the input text is too lengthy or vague.

Also, the model has a poor performance when it comes to long-term dependencies. Although LSTMs usually perform better than vanilla RNNs in the capture of long-range contexts, degradation while processing long sequences is their attribute, as well. This will usually lead to out-of-context summaries, which would accurately capture the beginning part of an article but miss or misrepresent parts near the end.

Review: shiv sena chief uddhav thackeray has slammed prime minister narendra modi for power and taking away the independence  
Original summary: sos sena slams pm modi for power eos

Predicted summary: to to to to to to to to to to to to eos to

*Output of reduced dataset is inefficient as the model is not well trained.*

Thus based on the error analysis, future prospects could include possibly employing transformer-based encoders and refining the attention mechanism or having named entity recognition (NER) as a pre-processing step.

#### 5.5 Summary of Findings

The summarization models framed under this project showcased the promise that deep learning holds in creating meaningful summaries of large text content. Among all the given architectures, the 2-layer LSTM with attention played well in all quantitative and qualitative metrics. Attention mechanisms performed a vital role in enhancing the decoder to pay attention to those input parts most relevant for more informative and coherent summaries of the output.

The improved summary quality as a result of different model variations was clearly reflected in the

quantitative metrics, especially the ROUGE scores. The qualitative analysis further established that the generated summaries were fluently constructed and concise, managing to capture the gist of articles for the most part. Though some challenges such as repetition, hallucinations, and details omitted still exist, the project has been a strong starting point for further exploration.

These findings emphasize the importance of architectural improvements and data preprocessing in NLP tasks. Next, it could incorporate some more advanced models like Transformers or fine-tune on larger diverse datasets for stepping deeper into performance.

## **6. CONCLUSION**

The aim of this project was to create and evaluate deep learning-based models for automatic text summarization using Long Short Term Memory (LSTM) networks. The investigation started with a simple LSTM model and eventually advanced towards more complex architectures such as stacked LSTMs and attention-based models. Series of experiments showed that once integrated with attention mechanisms, the generated summaries became significantly better both in terms of coherence and contextual correctness.

The dataset itself came from a set of news articles, thus allowing the models to learn a kind of headline-style summarization. Quantitative evaluation according to ROUGE scores and qualitative analysis of output summaries reflected that the attention-enhanced LSTM outperformed all in performance and fluency. All these were done on Google Colab's GPU-enabled environment and thus provided sufficient resources to experiment efficiently.

Although the results were promising, the project had certain limitations, such as small repetition in summaries, occasional information loss and no exact named entity. These are a few challenges to further research and development.

## **7. FUTURE SCOPE**

The future work perhaps needs to extend into more advanced models like Transformers or BERT-based summarizers that would prove to handle long-range distance dependencies more maturely along with pointer generator networks, coverage, and reinforcement learning for optimizing the inherited redundancy and factual accuracy.

For instance, the current scope of the dataset can be widened to different domains, which can possibly take in natural language understanding tasks such as named entity recognition (NER) for improved outputs from the summarization engine. Such an improvement could make the system applicable in the real world-not only in journalism but also in education and content curation.

## 8. REFERENCES

1. Sutskever, I., Vinyals, O., & Le, Q. V. (2014).  
*Sequence to Sequence Learning with Neural Networks*.  
Advances in Neural Information Processing Systems, 27.  
<https://arxiv.org/abs/1409.3215>
2. Bahdanau, D., Cho, K., & Bengio, Y. (2015).  
*Neural Machine Translation by Jointly Learning to Align and Translate*.  
International Conference on Learning Representations (ICLR).  
<https://arxiv.org/abs/1409.0473>
3. Rush, A. M., Chopra, S., & Weston, J. (2015).  
*A Neural Attention Model for Abstractive Sentence Summarization*.  
Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing (EMNLP).  
<https://arxiv.org/abs/1509.00685>
4. See, A., Liu, P. J., & Manning, C. D. (2017).  
*Get To The Point: Summarization with Pointer-Generator Networks*.  
Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics.  
<https://arxiv.org/abs/1704.04368>
5. Keras Documentation.  
*Keras: Deep Learning for Humans*.  
<https://keras.io>
6. Chollet, F. (2015).  
*Keras: The Python Deep Learning library*.  
GitHub Repository.  
<https://github.com/keras-team/keras>
7. TensorFlow Documentation.  
*TensorFlow: An End-to-End Open Source Machine Learning Platform*.  
<https://www.tensorflow.org>
8. Google Colaboratory.  
*Colab: Research and Education Tools*.  
<https://colab.research.google.com>
9. Lin, C. Y. (2004).  
*ROUGE: A Package for Automatic Evaluation of Summaries*.  
Proceedings of the ACL Workshop on Text Summarization Branches Out.
10. News Summary Dataset – Kaggle.  
*News Category Dataset*.  
<https://www.kaggle.com/datasets/sunnysai12345/news-summary>